

# Architecting Intermediate Layers for Efficient Composition in End-to-End Data Science Pipelines

Supun Abeysinghe



# Motivation

Typical end-to-end data science pipelines are created using multiple libraries/frameworks/systems

e.g. Pandas + PyTorch



DASK



Keras



PyTorch



TensorFlow


XGBoost



MODIN

# Motivation


These systems go to great lengths to optimize the performance of their specific workloads, however, **end-to-end performance is suboptimal**



End-to-end performance for a ETL + Deep Learning Workload

# Motivation

1. Data copying/transformation at system boundaries
2. Lack of global optimization



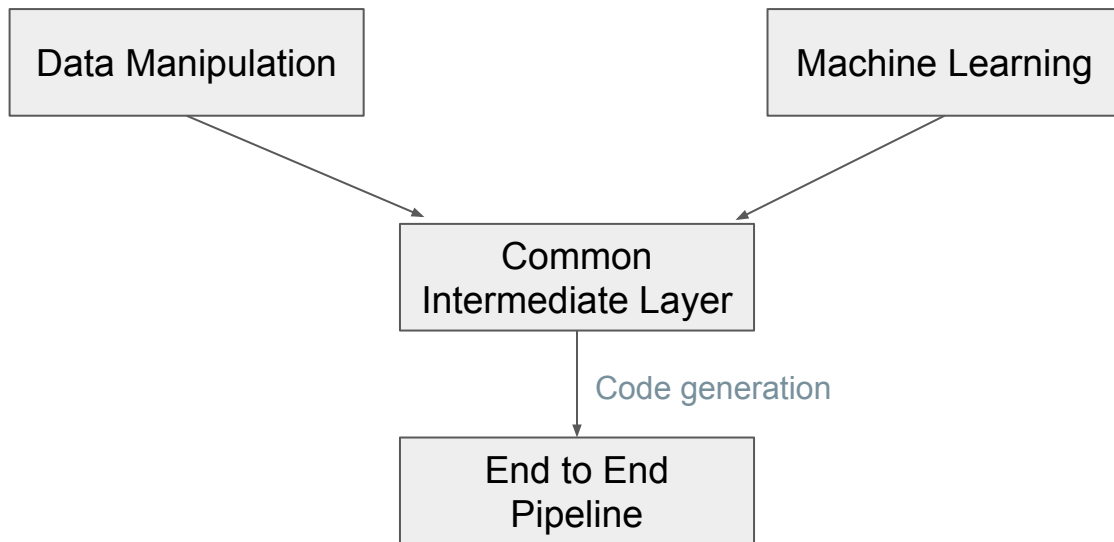
# Problem

How can we build efficient **end-to-end data science** pipelines that

1. avoid performance **overheads at subsystem boundaries**
2. do not preclude **global optimizations**
3. retain familiar **high-level interfaces** to the end user


# Intermediate Layers

A key idea to remove these bottlenecks is to reorganize existing frameworks around a **common intermediate layer**



# Flare

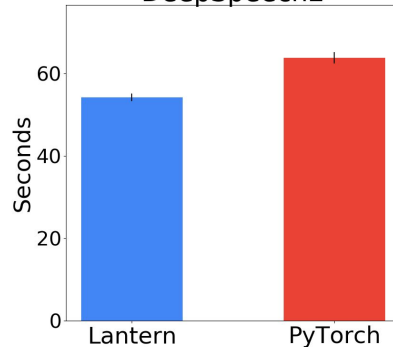
- Accelerator for Apache Spark
- Generates native code for optimized query plans from Apache Spark
- Leverages Lightweight Modular Staging (**LMS**)



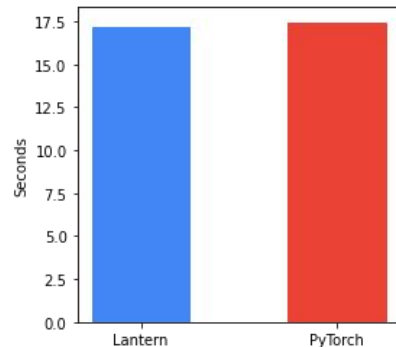
# Lantern

- Deep Learning framework
- Performance of “define-then-run”
- Expressiveness of “define-by-run”
- Leverages Lightweight Modular Staging (**LMS**) to generate low-level (C/CUDA) code


DeepSpeech2




Transformer



SqueezeNet



ResNet50





# Flern

**First** intermediate-layer integration between systems that achieves,

1. best-of-breed performance individually

competitive with best compiled query engines in **full relational benchmarks** (TPC-H)

competitive with Tensorflow, PyTorch in state-of-the-art machine learning models (**Transformer, DeepSpeech**)

2. competitive performance on end-to-end workloads

## Scala/Python Source

```
val df1 = spark.read(...)
val df2 = spark.read(...)
val df3 = spark.read(...)
```

```
val df4 = df1.withColumn("ratio", x /
y).withColumn(..)
```

```
query = spark.sql("SELECT .... FROM .... JOIN
... USING ... ORDER BY .....")
```

```
data = flare.executeQuery(query)
```

```
class NNModel {
  ....
}
```

```
opt = SGDOptimizer()
model = NNModel()
```

```
for (i <- 0 until nIter) {
  data.forEachBatch { input =>
    loss(model(input))
    opt.step()
  }
}
```

**Flare**

LMS

Low-level code  
(C)


**LMS**

Low-level code  
for the  
end-to-end task  
(C/CUDA)

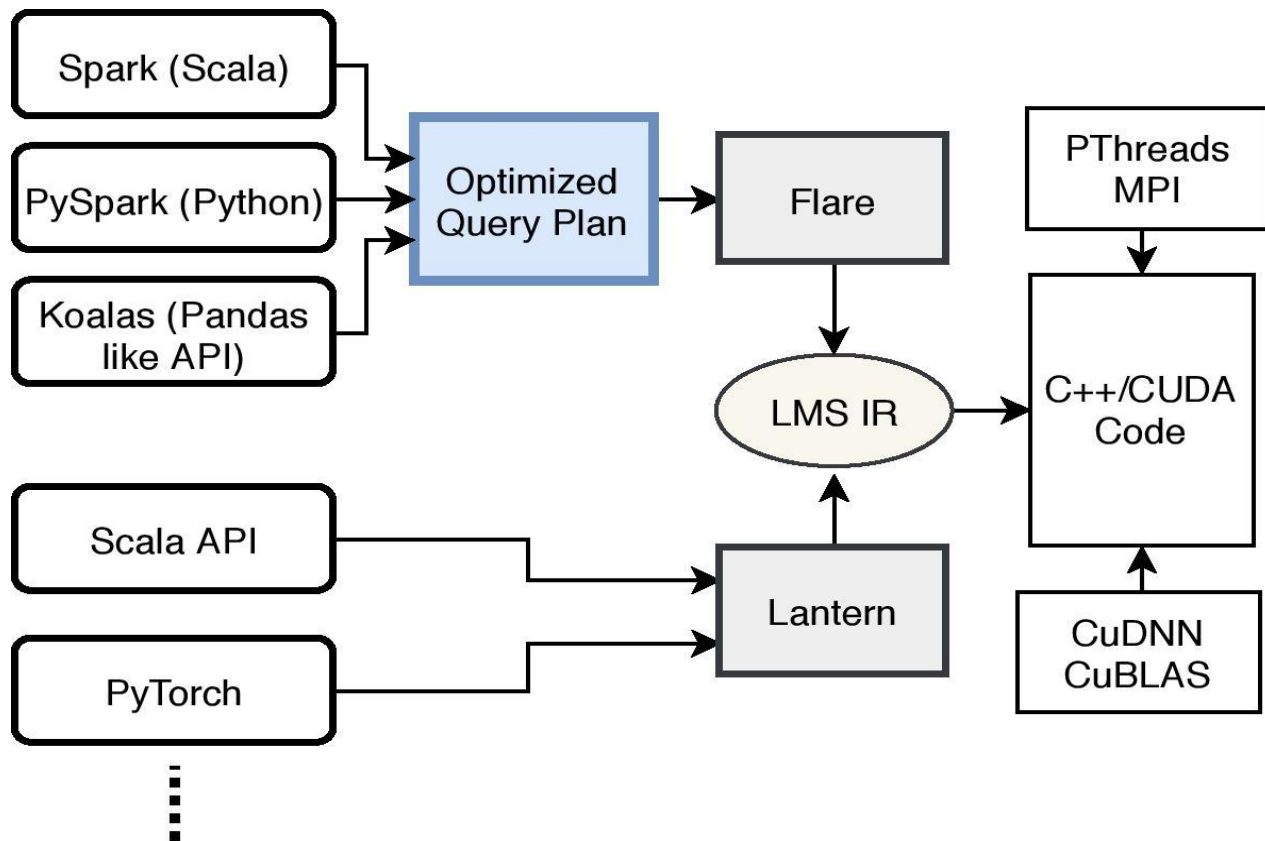
**Lantern**

LMS

Low-level code  
(C/CUDA)




# Overall Architecture




# Performance - End-to-end Training

4x speedup!



# Performance - End-to-end Inference

20x speedup!



# Future

Beyond single node; a lot of interesting challenges at the integration point!

Pipeline parallelism across end-to-end workloads

End to end GPU execution (+ generating specialized cuda kernels)

Traditional machine learning models (large scale logistic regression, gradient boosted trees, etc.)